

OpenOntos

From raw tables to a deployed warehouse — guided by AI agents.

This 10-step tour walks through the full OpenOntos workflow using the bundled **E-Commerce Data Warehouse** demo project. Each step shows a real screen capture from the live application and the AI Agent prompt that powers it.

01 Start: Project Workspace

03 Profile Source Data

05 Map Source → Target Attributes

07 Run & Schedule Pipelines

09 Monitor Pipeline Health

02 Connect Data Sources

04 Generate the Ontology

06 Generate ETL Code

08 Deploy to Target Platform

10 AI Agent — Everywhere

STEP 01 / 10


01

Start: Project Workspace

Pick a demo or create a new project

OpenOntos opens to your project workspace. The bundled **E-Commerce Data Warehouse** demo ships with 5 raw tables, a target ontology, and 10 mappings — perfect for a guided walkthrough. Click **+ New Project** to start from scratch.

AGENT *"Create a new project for retail sales analytics"*



STEP 02 / 10


02

Connect Data Sources

15+ supported source types — databases, warehouses, files, APIs

Connect PostgreSQL, MySQL, Snowflake, BigQuery, Databricks, Fabric, S3, REST APIs, and more. Each source is tested live and its tables are auto-discovered. Here the demo has a **PostgreSQL** database (4 tables) and a **REST API** (1 table) connected.

AGENT “Add a Snowflake source called Analytics with sample credentials”



STEP 03 / 10


03

Profile Source Data

AI-powered schema, quality, and relationship analysis

The Profile module ingests raw tables and runs **AI Analyze Schema** to detect column types, primary/foreign key candidates, anomalies, and join relationships. Anomalies (nulls, duplicates, outliers) are flagged inline — here `raw_customers` and `raw_order_items` have anomalies surfaced.

AGENT *“Profile all connected tables and explain any data quality issues”*



STEP 04 / 10


04

Generate the Ontology

Conceptual entities and relationships from your raw schema

The AI proposes a clean business ontology — **Customer** → **Order** → **Order Line Item** ← **Product** ← **Supplier** — with relationship types (Parent → Child, Contains, Related) and confidence scores per entity and edge. Edit, accept, or regenerate from the canvas.

AGENT “Suggest an ontology for an e-commerce order system”



STEP 05 / 10

05

Map Source → Target Attributes

AI-suggested column mappings with transformation expressions

Each ontology attribute is mapped from a source column with a generated transformation: CAST, TRIM, LOWER, COALESCE, etc. Confidence scores (90–99%) indicate how certain the AI is. Approve in bulk or refine per row. Here 6 mappings are approved, 4 pending.

AGENT “Suggest mappings for the Customer entity”

The screenshot displays the OpenOntos Mapping interface. The top navigation bar includes the OpenOntos logo, the current workspace 'E-Commerce Data Warehouse', and various utility buttons like 'DOCS', 'TOUR', 'RESET', 'REGISTER', and 'LOCAL v0.1.0-alpha'. Below the navigation bar, the 'Mapping' tab is selected, showing 'Source-to-target attribute mapping with AI suggestions'. The interface is divided into three main sections: 'Customer' (4 attributes), 'Order' (3 attributes), and 'Order Line Item' (1 attribute). Each section lists source columns, the AI-suggested transformation expression, and the target ontology attribute with a confidence score and a status (approved or pending).

Entity	Source Column	Transformation Expression	Target Attribute	Confidence Score	Status
Customer	raw_customers.customer_id	CAST(customer_id AS BIGINT)	Customer.customer_key	98%	Approved
	raw_customers.full_name	TRIM(UPPER(full_name))	Customer.customer_name	95%	Approved
	raw_customers.email	LOWER(TRIM(email))	Customer.email_address	93%	Approved
	raw_customers.segment	COALESCE(segment, 'standard')	Customer.customer_segment	90%	Approved
Order	raw_orders.order_id	CAST(order_id AS BIGINT)	Order.order_key	99%	Approved
	raw_orders.order_date	CAST(order_date AS TIMESTAMP)	Order.order_timestamp	96%	Approved
	raw_orders.total_amount	CAST(total_amount AS DECIMAL(12, ...))	Order.order_total	97%	Approved
Order Line Item					

STEP 06 / 10


06

Generate ETL Code

Bronze → Silver → Gold layered SQL for any target platform

OpenOntos generates production-ready SQL for **Snowflake**, **BigQuery**, **Redshift**, **Databricks**, and **Fabric** across Bronze (raw), Silver (cleaned), and Gold (aggregated) layers. Edit inline, copy, export, or execute directly on the warehouse.

AGENT *“Generate Snowflake ETL for all three layers”*



STEP 07 / 10


07

Run & Schedule Pipelines

Execute Bronze/Silver/Gold layers on demand or on schedule

Trigger individual layers from **Quick Run**, run the full pipeline end-to-end, or schedule recurring jobs (cron-style). Run history shows duration and row counts per execution. Switch target platforms with one click.

AGENT *"Schedule a nightly full refresh at 2 AM"*



STEP 08 / 10


08

Deploy to Target Platform

Configure connections and ship to production

Pick your warehouse — **Snowflake, Databricks, Microsoft Fabric, BigQuery, or Redshift** — and configure secure credentials. The **Readiness checklist** verifies schema profiled, ontology defined, mappings approved, ETL generated, and validations resolved before deployment.

AGENT *“Check deployment readiness and tell me what's blocking”*



STEP 09 / 10


09

Monitor Pipeline Health

Runs, validations, and AI health analysis

The Monitor view aggregates pipeline runs, errors, open validation issues, and active jobs. Click any run to drill into logs, row counts, and stack traces. **AI Health Analysis** summarizes failures and recommends fixes.

AGENT *"Why did the Gold Aggregation run fail and how do I fix it?"*



STEP 10 / 10


10

AI Agent — Everywhere

Context-aware copilot in every module

The **AI Agent panel** opens beside any module and adapts to its context — Profile Agent, Ontology Agent, Mapping Agent, Monitor Agent, and more. Quick prompts appear below the input. Each agent has full read/write access to your project state and can perform actions on your behalf.

AGENT “Show monitoring issues” · “Run validation checks” · “Summarize pipeline health”



Ready to build your own?

Open OpenOntos, click + **New Project**, and let the AI Agent guide you from raw data to deployed pipeline. Every module has its own contextual agent with quick-start prompts.

→ app.openontos.ai